

Independent single-layer adaptive predictors: same-state true-gradient references
Full batch; mean $\pm$ sample SD over 3 seeds; layer 8 reused

