



Predictor history transfers, but net model improvement remains unresolved

Gradient prediction: mean error -14.6%



One-update accuracy: carry > reset, near baseline



Primary model state 32 · 3 paired seeds · left: 4 held-out batches per seed · right: prespecified probe 0, 1,024 shared questions

Qwen2.5-0.5B · layer-8 rank-64 LoRA · full input prompts preserved · 2,048-token limit applies only to new output