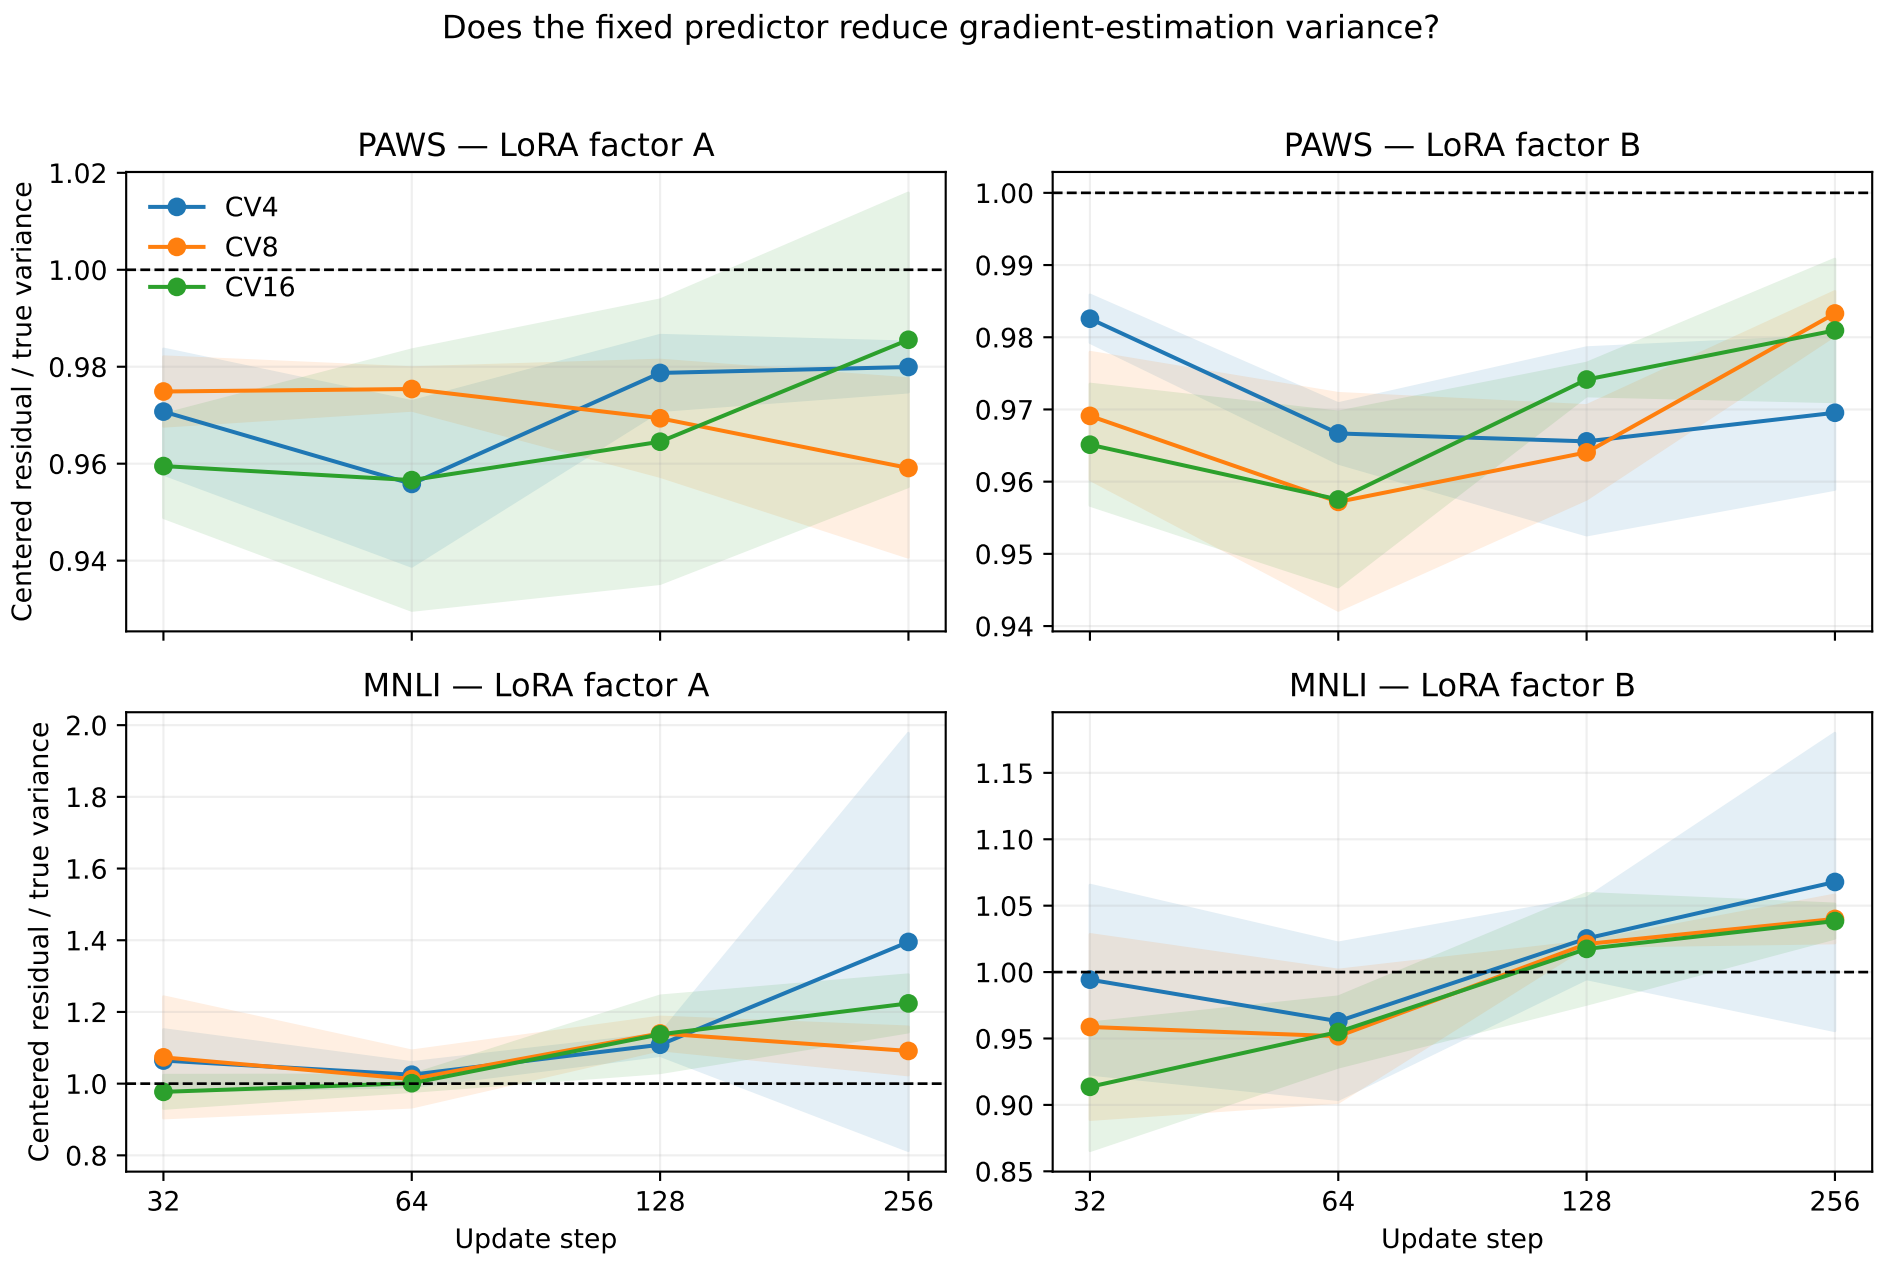
Does the fixed predictor reduce gradient-estimation variance?
PAWS — LoRA factor A
PAWS — LoRA factor B
MNLI — LoRA factor A
MNLI — LoRA factor B
CV4
CV8
CV16
Centered residual / true variance
Update step
16 held-out examples per checkpoint; descriptive only. Mean ± seed SD. Below 1 indicates variance reduction.