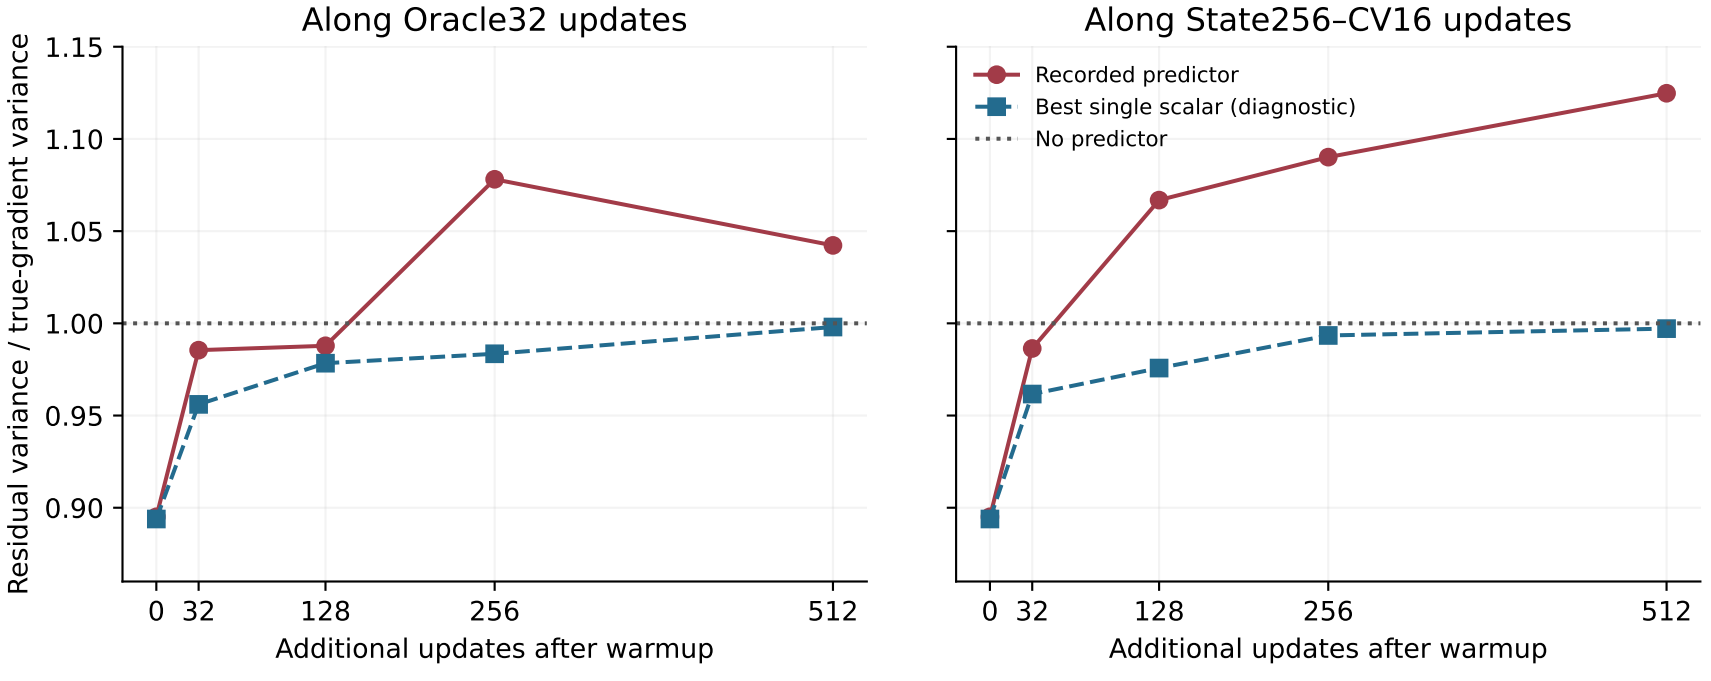

129 development examples; start 301; LR 0.001. Lower ratios are better; 1 is the no-predictor control.
The best scalar uses all true diagnostic gradients and is not a deployable calibration policy.
This fixed-state variance calculation is not a downstream-accuracy or trajectory guarantee.